



VIT[®]

Vellore Institute of Technology

(Deemed to be University under section 3 of UGC Act, 1956)

OPERATING SYSTEMS

DIGITAL ASSIGNMENT

STUDENT	AATMIKA DIXIT 22BPS1183
STUDENT	ASHISH 22BAI1062
STUDENT	MILIND RANJAN 22BPS1167
STUDENT	VEERAJ KUTE 22BPS1129
COURSE	BCSE303P
CLASS NO	CH2023240100703
FACULTY	VALLIDEVI K

Processing large amounts of data in an operating system involves several techniques and algorithms. Here are some key approaches and well-known algorithms:

- **Memory Management:** The OS manages memory efficiently to ensure that data can be loaded and processed. This includes techniques like paging and segmentation.
- **File Systems:** The OS uses file systems to organise and store large volumes of data. Popular file systems include NTFS, ext4, and HFS+.
- **I/O Scheduling:** Algorithms like the C-SCAN, FCFS, and SSTF are used to schedule disk I/O requests, optimising data access.
- **Buffering:** Buffering involves using a buffer to temporarily hold data, reducing the frequency of I/O operations. Techniques like double buffering and circular buffering are used.
- **Caching:** Caching algorithms like LRU (Least Recently Used) and LFU (Least Frequently Used) help improve data access times by keeping frequently used data in memory.
- **Data Compression:** OS can use compression algorithms like Lempel-Ziv-Welch (LZW) or Deflate to reduce the size of data for storage and transmission.
- **Parallel Processing:** Multithreading and multiprocessing techniques allow the OS to distribute data processing tasks across multiple cores or processors.
- **Virtual Memory:** Virtual memory techniques, including demand paging and page replacement algorithms like LRU and FIFO, are used to efficiently manage memory for large datasets.
- **Data Indexing:** For searching and retrieval of data, indexing algorithms like B-trees, Hashing, and various search algorithms are employed.
- **Distributed File Systems:** In a networked environment, distributed file systems like Hadoop HDFS and Google File System (GFS) are used to process and store large datasets across multiple machines.
- **Data Partitioning:** Large datasets can be divided into smaller partitions, and MapReduce-style algorithms are used for parallel processing, with frameworks like Hadoop.
- **Database Management Systems:** When dealing with large datasets in a database, DBMS systems like MySQL, PostgreSQL, and NoSQL databases are employed for efficient data retrieval and management.

- **Streaming Algorithms:** For processing data streams, algorithms like Count-Min Sketch and HyperLogLog are used to estimate statistics from large data streams in real-time.

Operating systems use scheduling algorithms to manage the execution of multiple processes or threads efficiently. The scheduling algorithm's primary goal is to allocate CPU time to different tasks or processes while maximising system throughput, fairness, and responsiveness. Here are some common scheduling algorithms used by operating systems:

- **First-Come, First-Served (FCFS):**

- - It schedules processes based on their arrival time. The process that arrives first gets executed first.
- - It is simple but can lead to poor utilisation of CPU time, especially if long processes arrive early (known as the "convoy effect").

- **Shortest Job First (SJF):**

- - Prioritises processes with the smallest CPU burst or execution time.
- - It minimises average waiting time but requires knowledge of each process's execution time, which may not always be available.

- **Round Robin (RR):**

- - Allocates fixed time slices (time quantum) to each process in a circular manner.
- - Provides fairness and prevents a single long-running process from hogging the CPU.
- - Shorter time slices lead to more frequent context switches.

- **Priority Scheduling:**

- - Assigns a priority to each process, and the CPU is allocated to the highest priority process.
- - Can be either preemptive (higher priority process can interrupt a lower one) or non-preemptive.

- **Multilevel Queue Scheduling:**

- - Divides the ready queue into multiple priority queues, each with its own scheduling algorithm.
- - Higher-priority queues receive more CPU time, and lower-priority queues are serviced only when the higher-priority queues are empty.

- **Multilevel Feedback Queue Scheduling:**

- - A variation of multilevel queue scheduling that allows processes to move between queues based on their behaviour (e.g., aging, I/O-bound vs. CPU-bound).
- - Provides adaptability and dynamic priority adjustments.

- **Lottery Scheduling:**

- - Allocates "lottery tickets" to processes, and a random drawing determines the process to execute.
- - Provides a way to implement probabilistic scheduling and fairness.

- **Completely Fair Scheduler (CFS):**

- - Used in Linux, it aims to divide CPU time fairly among processes, using a red-black tree to maintain the execution order.
- - Achieves fairness and responsiveness in dynamic workloads.

- **Earliest Deadline First (EDF):**

- - Used in real-time systems, processes are scheduled based on their deadlines.
- - Ensures that processes with the earliest deadlines get CPU time, critical for real-time applications.

Ubuntu, like many other Linux distributions, uses the Completely Fair Scheduler (CFS) as its default process scheduler. The CFS is part of the Linux kernel's scheduler and aims to provide a fair and efficient allocation of CPU time to processes. It uses a red-black tree to maintain the execution order and strives to divide CPU time fairly among processes, ensuring that no process hogs the CPU.

Simulation for CFS:

Code:

```
1. #include <stdio.h>
2. #include <stdlib.h>
3. #include <pthread.h>
4. #include <unistd.h>
5. #include <time.h>
6. #include <sched.h>
7.
8. #define NUM_THREADS 40
9. #define DATA_FILE "data.txt"
10. #define RESULT_FILE "results.txt"
11.
12. void *thread_function(void *arg) {
13.     int thread_id = *(int *)arg;
```

```

14.     int sum = 0;
15.
16.     // Open the data file for reading
17.     FILE *data_file = fopen(DATA_FILE, "r");
18.     if (data_file == NULL) {
19.         perror("Failed to open data file");
20.         pthread_exit(NULL);
21.     }
22.
23.     // Read integers from the data file and calculate the sum
24.     int num;
25.     while (fscanf(data_file, "%d", &num) == 1) {
26.         sum += num;
27.     }
28.
29.     fclose(data_file);
30.
31.     // Get the core number where the thread is running
32.     int core_num = sched_getcpu();
33.
34.     // Open the results file for writing
35.     FILE *results_file = fopen(RESULT_FILE, "a");
36.     if (results_file == NULL) {
37.         perror("Failed to open results file");
38.         pthread_exit(NULL);
39.     }
40.
41.     // Write the result, thread ID, and core number to the
    results file
42.     fprintf(results_file, "Thread %d (Core %d): Sum = %d\n",
    thread_id, core_num, sum);
43.     fclose(results_file);
44.
45.     pthread_exit(NULL);
46. }
47.
48. int main() {
49.     pthread_t threads[NUM_THREADS];
50.     int thread_ids[NUM_THREADS];
51.     clock_t start, end;
52.     double cpu_time_used;
53.
54.     start = clock();
55.
56.     // Open the results file and clear its contents
57.     FILE *results_file = fopen(RESULT_FILE, "w");
58.     if (results_file == NULL) {
59.         perror("Failed to open results file");
60.         return 1;
61.     }
62.     fclose(results_file);
63.

```

```

64.     for (int i = 0; i < NUM_THREADS; i++) {
65.         thread_ids[i] = i;
66.         int result = pthread_create(&threads[i], NULL,
thread_function, &thread_ids[i]);
67.         if (result) {
68.             perror("Thread creation failed");
69.             return 1;
70.         }
71.     }
72.
73.     for (int i = 0; i < NUM_THREADS; i++) {
74.         pthread_join(threads[i], NULL);
75.     }
76.
77.     end = clock();
78.     cpu_time_used = ((double)(end - start)) / CLOCKS_PER_SEC;
79.
80.     printf("All threads have finished. Total execution time: %f
seconds.\n", cpu_time_used);
81.
82.     return 0;
83. }

```

Terminal Window After Execution:

```

ashish_22bai1062@ubuntu:~$ cd OS_DA
ashish_22bai1062@ubuntu:~/OS_DA$ touch cfs.c
ashish_22bai1062@ubuntu:~/OS_DA$ touch data.c
ashish_22bai1062@ubuntu:~/OS_DA$ touch data.txt
ashish_22bai1062@ubuntu:~/OS_DA$ touch results.txt
ashish_22bai1062@ubuntu:~/OS_DA$ gcc cfs.c -o cfs -lpthread -lrt
cfs.c: In function 'thread_function':
cfs.c:32:20: warning: implicit declaration of function 'sched_getcpu'; did you mean
'sched_getparam'? [-Wimplicit-function-declaration]
   32 |         int core_num = sched_getcpu();
      |                        ^~~~~~
      |                        sched_getparam
ashish_22bai1062@ubuntu:~/OS_DA$ ./cfs
All threads have finished. Total execution time: 0.011944 seconds.
ashish_22bai1062@ubuntu:~/OS_DA$

```

I


Data File:

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47	48	49	50	51																					
52	53	54	55	56	57	58	59	60	61	62	63	64	65	66	67	68	69	70	71	72	73	74	75	76	77	78	79	80	81	82	83	84	85	86	87	88	89	90	91	92	93	94	95	96	97	98	99																								
100	101	102	103	104	105	106	107	108	109	110	111	112	113	114	115	116	117	118	119	120	121	122	123	124	125	126	127	128	129	130	131	132	133	134	135	136	137	138	139	140	141	142	143	144	145	146	147	148	149	150	151	152	153	154	155	156	157	158	159	160	161	162	163	164	165	166	167	168	169	170	171
172	173	174	175	176	177	178	179	180	181	182	183	184	185	186	187	188	189	190	191	192	193	194	195	196	197	198	199	200	201	202	203	204	205	206	207	208	209	210	211	212	213	214	215	216	217	218	219	220	221	222	223	224	225	226	227	228	229	230	231	232	233	234	235	236	237	238	239	240	241	242	243
244	245	246	247	248	249	250	251	252	253	254	255	256	257	258	259	260	261	262	263	264	265	266	267	268	269	270	271	272	273	274	275	276	277	278	279	280	281	282	283	284	285	286	287	288	289	290	291	292	293	294	295	296	297	298	299	300	301	302	303	304	305	306	307	308	309	310	311	312	313	314	315
316	317	318	319	320	321	322	323	324	325	326	327	328	329	330	331	332	333	334	335	336	337	338	339	340	341	342	343	344	345	346	347	348	349	350	351	352	353	354	355	356	357	358	359	360	361	362	363	364	365	366	367	368	369	370	371	372	373	374	375	376	377	378	379	380	381	382	383	384	385	386	387
388	389	390	391	392	393	394	395	396	397	398	399	400	401	402	403	404	405	406	407	408	409	410	411	412	413	414	415	416	417	418	419	420	421	422	423	424	425	426	427	428	429	430	431	432	433	434	435	436	437	438	439	440	441	442	443	444	445	446	447	448	449	450	451	452	453	454	455	456	457	458	459
460	461	462	463	464	465	466	467	468	469	470	471	472	473	474	475	476	477	478	479	480	481	482	483	484	485	486	487	488	489	490	491	492	493	494	495	496	497	498	499	500	501	502	503	504	505	506	507	508	509	510	511	512	513	514	515	516	517	518	519	520	521	522	523	524	525	526	527	528	529	530	531
532	533	534	535	536	537	538	539	540	541	542	543	544	545	546	547	548	549	550	551	552	553	554	555	556	557	558	559	560	561	562	563	564	565	566	567	568	569	570	571	572	573	574	575	576	577	578	579	580	581	582	583	584	585	586	587	588	589	590	591	592	593	594	595	596	597	598	599	600	601	602	603
604	605	606	607	608	609	610	611	612	613	614	615	616	617	618	619	620	621	622	623	624	625	626	627	628	629	630	631	632	633	634	635	636	637	638	639	640	641	642	643	644	645	646	647	648	649	650	651	652	653	654	655	656	657	658	659	660	661	662	663	664	665	666	667	668	669	670	671	672	673	674	675
676	677	678	679	680	681	682	683	684	685	686	687	688	689	690	691	692	693	694	695	696	697	698	699	700	701	702	703	704	705	706	707	708	709	710	711	712	713	714	715	716	717	718	719	720	721	722	723	724	725	726	727	728	729	730	731	732	733	734	735	736	737	738	739	740	741	742	743	744	745	746	747
748	749	750	751	752	753	754	755	756	757	758	759	760	761	762	763	764	765	766	767	768	769	770	771	772	773	774	775	776	777	778	779	780	781	782	783	784	785	786	787	788	789	790	791	792	793	794	795	796	797	798	799	800	801	802	803	804	805	806	807	808	809	810	811	812	813	814	815	816	817	818	819
820	821	822	823	824	825	826	827	828	829	830	831	832	833	834	835	836	837	838	839	840	841	842	843	844	845	846	847	848	849	850	851	852	853	854	855	856	857	858	859	860	861	862	863	864	865	866	867	868	869	870	871	872	873	874	875	876	877	878	879	880	881	882	883	884	885	886	887	888	889	890	891
892	893	894	895	896	897	898	899	900	901	902	903	904	905	906	907	908	909	910	911	912	913	914	915	916	917	918	919	920	921	922	923	924	925	926	927	928	929	930	931	932	933	934	935	936	937	938	939	940	941	942	943	944	945	946	947	948	949	950	951	952	953	954	955	956	957	958	959	960	961	962	963
964	965	966	967	968	969	970	971	972	973	974	975	976	977	978	979	980	981	982	983	984	985	986	987	988	989	990	991	992	993	994	995	996	997	998	999																																				

Results File:

```
Thread 2 (Core 0): Sum = 499500
Thread 0 (Core 2): Sum = 499500
Thread 1 (Core 0): Sum = 499500
Thread 3 (Core 3): Sum = 499500
Thread 5 (Core 3): Sum = 499500
Thread 4 (Core 0): Sum = 499500
Thread 7 (Core 1): Sum = 499500
Thread 8 (Core 2): Sum = 499500
Thread 6 (Core 3): Sum = 499500
Thread 9 (Core 2): Sum = 499500
Thread 12 (Core 3): Sum = 499500
Thread 11 (Core 2): Sum = 499500
Thread 16 (Core 3): Sum = 499500
Thread 13 (Core 3): Sum = 499500
Thread 15 (Core 2): Sum = 499500
Thread 19 (Core 3): Sum = 499500
Thread 18 (Core 2): Sum = 499500
Thread 14 (Core 0): Sum = 499500
Thread 20 (Core 2): Sum = 499500
Thread 10 (Core 0): Sum = 499500
Thread 17 (Core 0): Sum = 499500
Thread 22 (Core 2): Sum = 499500
Thread 21 (Core 0): Sum = 499500
Thread 23 (Core 3): Sum = 499500
Thread 24 (Core 0): Sum = 499500
Thread 26 (Core 0): Sum = 499500
Thread 29 (Core 1): Sum = 499500
Thread 27 (Core 2): Sum = 499500
Thread 25 (Core 1): Sum = 499500
Thread 28 (Core 0): Sum = 499500
Thread 30 (Core 0): Sum = 499500
Thread 32 (Core 1): Sum = 499500
Thread 33 (Core 0): Sum = 499500
Thread 38 (Core 0): Sum = 499500
Thread 37 (Core 1): Sum = 499500
Thread 31 (Core 3): Sum = 499500
Thread 35 (Core 3): Sum = 499500
Thread 39 (Core 2): Sum = 499500
Thread 34 (Core 3): Sum = 499500
Thread 36 (Core 2): Sum = 499500
```

File System View:



A Real-Time Operating System (RTOS) is a specialized operating system designed for systems and applications that require precise and deterministic timing and rapid response to external events. Unlike general-purpose operating systems (such as Windows, Linux, or macOS), which are optimized for multitasking and may prioritize tasks based on various factors, RTOS is focused on meeting stringent timing requirements and ensuring that tasks are executed within specific time constraints. Here are some key characteristics and features of real-time operating systems:

1. ****Deterministic Behavior:**** RTOS systems are known for their deterministic behavior. They guarantee that tasks will be executed within predetermined time frames, making them suitable for applications where timing is critical. This predictability is crucial in fields like aerospace, automotive, industrial automation, and medical devices.
2. ****Prioritization:**** Real-time tasks are assigned priorities, and the scheduler ensures that higher-priority tasks are executed before lower-priority tasks. This priority-based scheduling helps meet timing deadlines.

3. ****Hard Real-Time vs. Soft Real-Time:**** RTOS can be categorized into hard real-time and soft real-time systems. In a hard real-time system, missing a deadline can have catastrophic consequences, while in a soft real-time system, occasional missed deadlines might be acceptable. Hard real-time systems demand more stringent timing guarantees.

4. ****Task Synchronization and Communication:**** RTOS provides mechanisms for tasks to synchronize and communicate with each other. Common synchronization tools include semaphores, mutexes, and message queues, which ensure that tasks do not interfere with each other.

5. ****Interrupt Handling:**** RTOS is equipped to efficiently handle hardware interrupts and manage real-time tasks in response to these events. This is crucial for systems where immediate response to external events is vital.

6. ****Minimal Overhead:**** RTOS is designed to have minimal overhead. This means that the operating system itself is lightweight and doesn't introduce significant delays or uncertainty into task execution.

7. ****Resource Management:**** Real-time operating systems manage resources efficiently. This includes memory allocation, CPU utilization, and other system resources to ensure that they are allocated optimally.

8. ****Embedded Systems:**** RTOS is commonly used in embedded systems, such as microcontrollers, where it is crucial to meet strict timing requirements in applications like automotive control systems, medical devices, robotics, and communication equipment.

9. ****Task Constraints:**** RTOS allows you to define timing constraints for tasks, such as maximum execution time, periodicity, and deadlines, to ensure that the system operates within the required parameters.

10. **Safety-Critical Applications:** Real-time operating systems are often used in safety-critical applications where reliability and predictability are paramount. This includes systems in the aerospace industry, nuclear reactors, and medical devices.

Some well-known real-time operating systems include:

- **FreeRTOS:** An open-source RTOS with a large user base and a wide range of supported microcontrollers.
- **VxWorks:** A commercial real-time operating system used in various critical applications, including aerospace and automotive.
- **QNX:** Known for its reliability and used in applications like automotive infotainment systems and medical devices.
- **RTOS-32:** An RTOS designed for Windows systems, providing real-time capabilities on Windows platforms.

RTOS plays a crucial role in ensuring that critical systems operate with high reliability and predictability. When choosing an RTOS for a specific application, factors like timing requirements, hardware compatibility, and licensing costs should be carefully considered.

CODE:

```
1. #include <stdio.h>
2. #include "FreeRTOS.h"
3. #include "task.h"
4. #include "semphr.h"
5. #include "esp_timer.h"
6.
7. #define NUM_TASKS 200
8. #define NUM_CORES 4
9. #define DATA_FILE "data.txt"
10. #define RESULT_FILE "results.txt"
11.
12. // Define a mutex to protect the file access
13. SemaphoreHandle_t fileMutex;
14.
```

```

15. void thread_function(void *pvParameters) {
16.     int task_id = (int)pvParameters;
17.     int sum = 0;
18.
19.     // Record the start time
20.     int64_t start_time = esp_timer_get_time();
21.
22.     // Open the data file for reading
23.     FILE *data_file = fopen(DATA_FILE, "r");
24.     if (data_file == NULL) {
25.         perror("Failed to open data file");
26.         vTaskDelete(NULL);
27.     }
28.
29.     // Read integers from the data file and calculate
    the sum
30.     int num;
31.     while (fscanf(data_file, "%d", &num) == 1) {
32.         sum += num;
33.     }
34.
35.     fclose(data_file);
36.
37.     // Record the end time
38.     int64_t end_time = esp_timer_get_time();
39.
40.     // Get the core number where the task is running
41.     int core_num = xPortGetCoreID();
42.
43.     // Take the mutex to protect file access
44.     if (xSemaphoreTake(fileMutex, portMAX_DELAY) ==
pdTRUE) {
45.         // Open the results file for writing
46.         FILE *results_file = fopen(RESULT_FILE, "a");
47.         if (results_file == NULL) {
48.             perror("Failed to open results file");
49.             vTaskDelete(NULL);
50.         }
51.
52.         // Write the result, task ID, core number, and
        execution time to the results file
53.         fprintf(results_file, "Task %d (Core %d): Sum =
%d, Start Time = %lld, End Time = %lld\n", task_id,
        core_num, sum, start_time, end_time);
54.         fclose(results_file);

```

```

55.
56.         // Release the mutex
57.         xSemaphoreGive(fileMutex);
58.     }
59.
60.     vTaskDelete(NULL);
61. }
62.
63. int main() {
64.     // Initialize the mutex for file access
65.     fileMutex = xSemaphoreCreateMutex();
66.
67.     if (fileMutex == NULL) {
68.         perror("Failed to create mutex");
69.         return 1;
70.     }
71.
72.     // Create FreeRTOS tasks on multiple cores
73.     TaskHandle_t tasks[NUM_TASKS];
74.     for (int i = 0; i < NUM_TASKS; i++) {
75.         int task_id = i;
76.         // Create tasks on different cores (cyclic
assignment)
77.         int core_num = task_id % NUM_CORES;
78.         xTaskCreatePinnedToCore(thread_function,
"Task", configMINIMAL_STACK_SIZE, (void *)task_id,
tskIDLE_PRIORITY + 1, &tasks[i], core_num);
79.     }
80.
81.     // Start the FreeRTOS scheduler
82.     vTaskStartScheduler();
83.
84.     return 0; // Should never reach here
85. }

```

Thank You!!